

Assignment 4: Fact-checking Outputs from ChatGPT

Deadline: week 14

Github invite link: <https://classroom.github.com/a/Un454G1q>

Goals The primary goal with this assignment is to give you hands-on experience analyzing outputs from large language models (LLMs) like ChatGPT. You will understand what sorts of non-factual outputs LLMs can generate, and what is involved in the process of verifying those outputs against Wikipedia. In addition, you will conduct some error analysis of the mistakes your fact-checking model makes.

Dataset and Code

Please use up-to-date versions of Python and PyTorch for this assignment. In addition, you will need to install the `transformers` package from Hugging Face as well as a few dependencies.

We have included a `requirements.txt` file. You can install the dependencies here by running:

```
pip install -r requirements.txt
```

Note that you can do this inside a virtual environment.

Data: The dataset and overall setup for this project is taken from FActScore (Min et al., 2023).¹ FActScore investigates the problem of detecting errors in biographies of people generated by ChatGPT. The FActScore paper explores methods to decompose claims into facts, then searches over Wikipedia to find articles that (potentially) support those facts. They conduct human annotation of these steps to understand the performance of systems in different parts of this pipeline (decomposition and fact-checking). The decomposition and retrieval have been done for you.

The **facts** you are given are propositions extracted from ChatGPT outputs. These are simple natural language sentences. Although FActScore explores different ways of producing these facts from biographies, we use their *human-annotated* facts for our analysis, to remove any potential errors in this step. These facts are then validated against **passages** that are retrieved by Wikipedia. Each fact is hand-labeled by humans with either “S” for supported, “NS” for not supported, or “IR” for irrelevant. **However, we will only predict S vs. NS to simplify the task.**

Passages have been retrieved using BM25, a sparse retrieval model (i.e., one that does not use neural networks). The FActScore paper also discusses experiments with a dense retriever, but passages from this model only give slightly better performance.

In the provided Python code, `FactExample` stores the examples themselves. These consist of a fact (a string), passages (a list of dicts consisting of a title and text), and the label. Note that the passages themselves are not tokenized or preprocessed in any way.

The raw fact data is located at `data/dev_labeled_ChatGPT.jsonl` and look like this (note: this example is not in the dataset given to you):

```
{ "input": "Question: Tell me a bio of Lanny Flaherty.",
  "output": "Lanny Flaherty is an American actor born on December 18, 1949, in Pensacola, Florida. He has appeared in numerous films, television shows, and theater productions throughout his career, which began in the late 1970s.",
  "label": "S",
  "passages": [
    { "title": "Lanny Flaherty", "text": "Lanny Flaherty is an American actor born on December 18, 1949, in Pensacola, Florida. He has appeared in numerous films, television shows, and theater productions throughout his career, which began in the late 1970s." }
  ]
}
```

¹Original GitHub: <https://github.com/shmsw25/FActScore>

```

        Pensacola, Florida.",
    "is-relevant": true,
    "human-atomic-facts": [
        {"text": "Lanny Flaherty is an American.", "label": "S"},
        {"text": "Lanny Flaherty is an actor.", "label": "S"},
        {"text": "Lanny Flaherty was born on December 18, 1949.", "label": "NS"} [...]]

```

Note that the output field is what ChatGPT generated about this person, which may or may not be correct. The human-atomic-facts are annotated facts with labels indicating whether each part of the output is true or not, according to humans who were able to consult Wikipedia.

The retrieved passages are in `passages_bm25_ChatGPT_humfacts.jsonl` and they look like this. Note that only a single passage was retrieved in this case, but the passages you have are generally more numerous. This is information from Wikipedia, so is likely to be more reliable than what the model generated.

```

{"name": "Lanny Flaherty",
 "sent": "Lanny Flaherty is an American.",
 "passages": [{"title": "Lanny Flaherty",
  "text": "<s>Lanny Flaherty Lanny Flaherty (born July 27, 1942) is an
American actor.</s><s>Career. He has given his most memorable performances
in \"Lonesome Dove\", \"Natural Born Killers\", \"\" and \"Signs\". Flaherty
attended University of Southern Mississippi after high school. He also
had a brief role in \"Men in Black 3\", and appeared as Jack Crow in Jim
Mickles 2014 adaptation of \"Cold in July\". Other film appearances include
\"Winter People\", \"Millers Crossing\", \"Blood In Blood Out\", \"Tom and
Huck\" and \"Home Fries\" while television roles include guest appearances
on \"The Equalizer\", \"New York News\" and \"White Collar\" as well as a 2
episode stint on \"The Education of Max Bickford\" as Whammo.</s><s>Personal
life. Flaherty resides in New York City.</s>\"}]]}

```

Framework code: The framework code you are given consists of two files:

1. `factchecking_main.py`: This loads the data, instantiates the `FactChecker` model to use, executes it on the data, and reports accuracy, similar to the driver classes in previous projects. You will not be modifying this file.
2. `factchecker.py`: This file contains `FactExample`, `EntailmentModel`, and the `FactChecker` classes you will implement. This is similar to `models.py` in previous assignments.

Getting started: Run:

```

python factchecking_main.py --mode random
python factchecking_main.py --mode always_entail

```

This will run two low-performing baselines and print results. You will implement stronger `FactChecker` models that will perform better!

Part 1: Word Overlap (35 points)

The first method we will explore for fact-checking is bag-of-words overlap. A supported fact might be expected to have high unigram overlap with a passage that supports it, and low overlap with passages that don't support it. This sounds like a somewhat naive baseline, but it's one that might work pretty well!

You should implement a method that predicts supported vs. not supported based on this criterion. You will have to decide on several details of your implementation. How do you want to handle tokenization? Do you want to stem? Do you want to remove stopwords? The decisions you are making here look similar to those in the bag-of-words method in Assignment 1, but what works best may be different!

Your method should eventually compute a real-valued score that you use to make your classification decision. This could be a cosine similarity of tf-idf vectors, Jaccard similarity of word sets, metrics based on precision/recall, existing lexical metrics like ROUGE or BLEU, or anything similar that computes a discrete (non-neural/embedding-based) word overlap score.

Your method should get at least 75% accuracy on the dataset to receive full credit. Try tuning the preprocessing method as well as changing the classification threshold to achieve the best performance. You can tune this threshold on the development set to optimize for accuracy. It would be best practice to tune this on a separate set, but we won't have separate dev and test sets here.

Part 2: Textual Entailment (35 points)

Second, you will explore a method that goes beyond the surface form of the words. The most appropriate type of system explored in prior work is textual entailment. These systems should theoretically be able to decide whether a fact (the “hypothesis”) is logically entailed by a source document (the “premise”).

For this part, you will use a pre-trained model, specifically a DeBERTa-v3 base model. This is a pre-trained instance of DeBERTa (He et al., 2020) fine-tuned on the MNLI (Williams et al., 2018), FEVER (Thorne et al., 2018), and ANLI (Nie et al., 2020). This model can take a premise-hypothesis pair and return a decision: entailment, neutral, or contradiction. Your task is to use this information to determine supported vs. not supported for the fact.

You can run

```
python factchecking_main.py --mode entailment --cuda
```

with the optional `--cuda` flag to do the entailment model queries on a GPU. You should do this only if you have Pytorch with CUDA installed and a GPU enabled.

Using the entailment model: The model itself is loaded in `factchecking_main.py`. In `factcheck.py`, the `EntailmentModel` class includes most of the boilerplate code needed to do inference with this model. However, we stop short of mapping logits to an actual prediction. You will have to decide what strategy makes the most sense for extracting a binary S/NS decision from a three-class entailment/neutral/contradiction decision. You can either use the discrete entailment decision, or you can derive a decision based on a threshold of one of the entailment/neutral/contradiction probabilities (or some combination thereof).

Entailment models are typically designed to take sentences as input, not entire passages.² Therefore, you will have to do some kind of sentence splitting and likely some “cleaning” as well, due to noise in the data introduced by Wikipedia. You should then loop over the sentences and systematically compare the fact to each sentence in the passages, then take the “max” over the passages.

If your implementation uses the entailment classifier discretely (e.g., taking the returned entailment label), then the sentence can be considered entailed if it is entailed by *any* part of any of the passages. If your implementation uses entailment scores and sets a threshold, then the score for that fact should be the

²The ANLI dataset does contain paragraphs as passages. However, in our experimentation, we did not find that feeding whole-passage premises into the model worked well. You are free to experiment with it, though!

max over the entailment scores returned by running on any example in the passage. This resembles the method presented in Laban et al. (2022).

Your method should get at least 83% accuracy on the dataset to receive full credit.

Optimization: Running the entailment model on all of the data may be slow. To speed things up, you can implement pruning based on the word overlap method. If there is very low word overlap, then entailment is extremely unlikely. You can discard examples that fail to meet a low word overlap threshold, then run entailment on the remaining instances.

Aiming to complete in around 10 minutes on your machine is a good target. Second, your code must run without causing out-of-memory (OOM) issues. We included a code snippet to remove variables that are no longer needed and explicitly delete large objects using the `del` function and `gc.collect()` function.

Part 3: Error Analysis (15 points)

Part 3 and Part 4 involve a written submission. These parts should be submitted together. See the end of the assignment PDF for submission instructions.

Finally, you should conduct some error analysis on the results of the **entailment** model (your model from Part 2).³

There are two broad types of errors you should investigate. **False positives** are cases where the model predicted “supported” but the ground truth label is “not supported.” **False negatives** are the opposite, where the model predicted “not supported” but the ground truth label is “supported.”

You should manually examine **10 errors** of each type (10 false positives and 10 false negatives), or as many errors as your system makes if it does not make 10. You should categorize them into a set of a few **fine-grained categories** that you see as suitable. For example, one kind of fine-grained error could be “the gold standard is wrong and the model is actually correct”. You should aim for 2-4 fine-grained error categories collectively between false positives and false negatives; you do not need to come up with many categories, but it’s important to analyze the data accurately. **These categories should not just be “false positive” or “false negative.”**

In your writeup, include:

1. Brief definitions of the fine-grained error categories you used in your analysis (1-2 sentences per category)
2. Aggregate statistics about the errors you examined: how many errors were of each fine-grained type?
3. **Three** examples discussed in more detail. For each of these, include: (a) the text of the example itself; (b) the ground truth label; (c) your model’s predicted label; (d) the fine-grained error type you gave it; (e) 1-3 sentences describing why you think the fine-grained error type label applies.

“Average” solutions on this part will receive full credit. It benefits you to be thorough for your own learning and as preparation for the final project. But you should not feel like you need to pad your analysis with lots of additional words to try to be “thorough.” Instead, aim to be precise and descriptive without being verbose.

³If you are not able to get Part 2 working, you are allowed to use the model from Part 1.

Part 4: Error Analysis of LLM Outputs (15 points)

Part 3 and Part 4 involve a written submission. These parts should be submitted together. See the end of the assignment PDF for submission instructions.

Finally, we want you to explore use cases around AI assistants. **You can opt-in for your responses to this part to be included as part of a research study; more details will follow.**

You will be investigating how LLMs perform and aspects of evaluating them across two realistic tasks. You will give LLMs instructions to do various tasks that you might use AI assistants for. For evaluation, we want you to analyze the aspects you use to see if the LLM response satisfied you or not. An aspect describes a dimension, such as factuality, structure, safety, logic, flow and more.

For instance, you might prompt an AI Assistant with the following: “Write me a 5-day travel itinerary for Austin, Texas around Austin City Limits. This is for a blog post.” Some of the aspects I might use for evaluation can be: (1) Correctness: the response correctly talks about things in Austin, Texas; (2) Engagingness: the response is engaging, has personal anecdotes and persuades the user to visit Austin, Texas; (3) Coherence: the response is well structured with good topical grouping of things

Your task: Please come up with **two different prompts** which might be useful to you and which also emphasize a balance of structural, tone and content constraints. Your prompts should follow the following constraints:

1. They should require a longer form answer (at least 1 paragraph) and shouldn’t simply be knowledge based (e.g.: when was Abraham Lincoln born?)
2. They should be something that you can evaluate for utility (e.g., writing a literature review, writing a tweet thread for your paper based on abstract and conclusion, writing an essay on a topic etc). We want you to be able to critique these prompts.
3. You should be comfortable critiquing the output to the question (don’t ask about quantum physics if you don’t know about it)

Then, get an answer from **two** of the following models: ChatGPT (www.chatgpt.com), Meta AI (www.meta.ai), or Gemini (<https://gemini.google.com/>). You may use any version of these models you like, either free or paid.

Finally, evaluate the response. Your evaluation should be done by you and not use AI assistants. You should evaluate **at least 3** aspects. For each response, what you submit should look like:

Prompt: [prompt]

Response (including model name): [response]

1. [Aspect name: your assessment and reasoning]
2. [Aspect name: your assessment and reasoning]
3. [Aspect name: your assessment and reasoning]

E.g., “Engagingness: the response is engaging and the personal anecdotes persuade me to visit Austin”

You will evaluate a total of 4 responses (2 prompts x 2 models) in this format with at least 3 aspects each. Each (prompt, model) combination can use distinct aspects, so you could have 12 (or more), or as few as 3 distinct aspects.

You should find a flaw in at least 2 of the responses. To achieve this, you can use one of the following strategies: (1) Ask about more obscure knowledge; (2) Place more constraints on the response; (3) Ask for more interesting output formats (a LinkedIn post, a TED talk, etc.).

Ideally your prompts should be as creative and distinct as you can make them: if something comes to mind from your work or a particular life situation, that's great. **However, please don't add any personal identifiable information that you are not comfortable with the course staff potentially seeing as they grade the assignments.**

Here are some examples of prompts you can use (**do not use these prompts directly**):

1. Write me a LinkedIn post talking about my promotion, add placeholders for any spot you need me to add personal details
2. Write a blog post/LinkedIn post/TED talk about how to be successful as a freelance ML engineer in the United States.
3. Write me a road trip itinerary for New Zealand's south island in the form of a blog. For any link or images, add a placeholder which I can add to later
4. Help me craft a statement of purpose for my PhD/Master's applications. I am [...]. Add placeholders for any information you need me to add later
5. I am organizing an NLP reading group on natural language generation evaluation with 6 computer science undergraduate students. How should I format the reading group and how do I organize it?
6. Help me write a twitter thread from the following abstract and conclusion. Add placeholders for where you think I need to add results or anything significant from the main paper. Abstract: [abstract] Conclusion: [conclusion]
7. I want to write an email to a friend who is afraid of public speaking, encouraging them to speak at a workshop I am organizing. Write me a draft of the email, helping me convince the friend.

Here are some examples of aspects you can draw on for your evaluation, but we encourage you to go beyond them!

1. Factuality: is the response factually correct?
2. Adherence to instruction: did the response follow your instructions
3. Tone: is the response in the tone you intended? Eg: casual, formal, sarcastic etc
4. Delivery: if you expected the response to be engaging or persuasive, is the model able to generate that?
5. Empathy: does the response contain the right level and style of empathy, is not condescending, and not too pleasing
6. Structure : is the response structured/organized as you intended (bullet, topical organization)
7. Reasoning: do ideas in the generation logically flow and build upon each other? (could be things like unsupported arguments, emotional manipulations, logic gaps – any fallacies that the generated output makes)
8. Completeness: right level of detail/length
9. Fluency: is the response fluent in English?
10. Coherence: is the response easy to follow and is everything flowing smoothly with things logically arranged and connected
11. Readability

Deliverables and Submission

Written Submission: You will upload your written submission as either a .doc file or a .pdf on **Github**.

Code Submission: You will upload your code for the first two parts on Github in a single file.

Make sure that the following commands work before you submit:

```
python factchecking_main.py --mode word_overlap
python factchecking_main.py --mode entailment
```

These commands should run without error and train in the allotted time limits.

References

- Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. 2020. DeBERTa: Decoding-enhanced BERT with Disentangled Attention. *ArXiv*, abs/2006.03654.
- Philippe Laban, Tobias Schnabel, Paul N. Bennett, and Marti A. Hearst. 2022. SummaC: Re-visiting NLI-based models for inconsistency detection in summarization. *Transactions of the Association for Computational Linguistics*, 10:163–177.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation. *arXiv 2305.14251*.
- Yixin Nie, Adina Williams, Emily Dinan, Mohit Bansal, Jason Weston, and Douwe Kiela. 2020. Adversarial NLI: A new benchmark for natural language understanding. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4885–4901, Online, July. Association for Computational Linguistics.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. 2018. FEVER: a large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 809–819, New Orleans, Louisiana, June. Association for Computational Linguistics.
- Adina Williams, Nikita Nangia, and Samuel Bowman. 2018. A Broad-Coverage Challenge Corpus for Sentence Understanding through Inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*.